

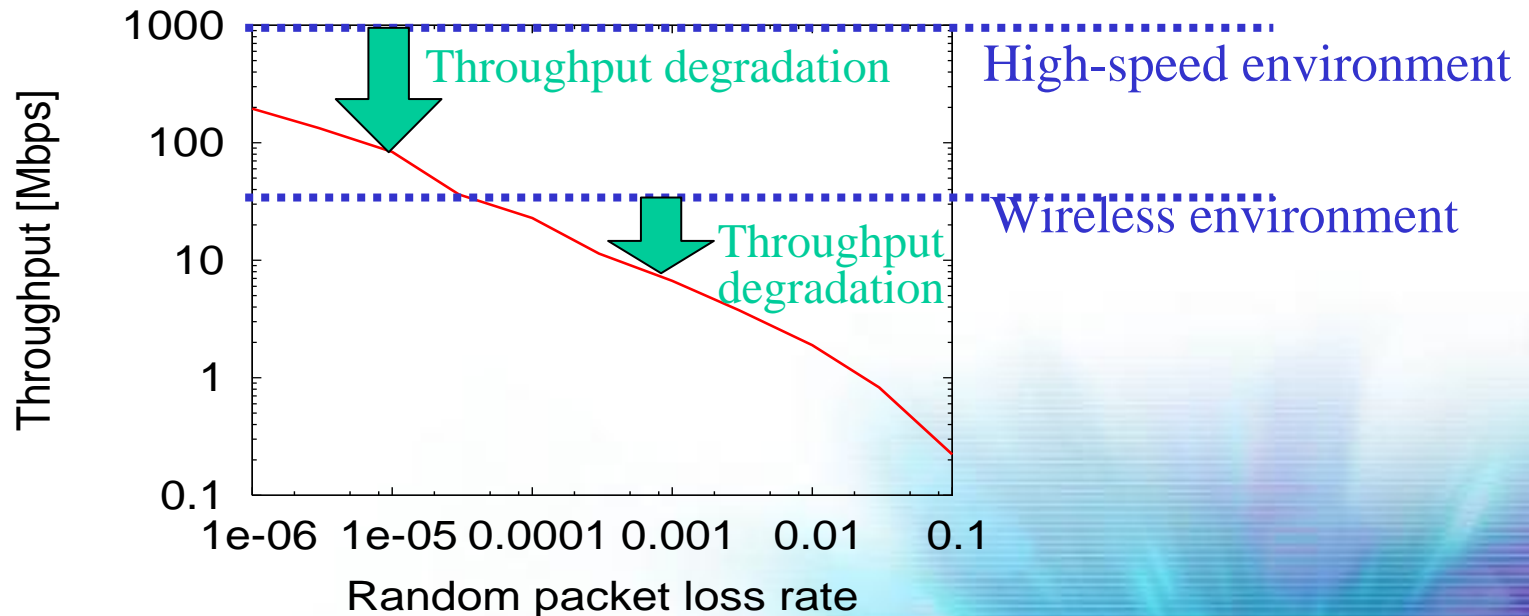
TCP-AdaptiveReno: Improving Efficiency-Friendliness Tradeoffs of TCP Congestion Control Algorithm

HIDEyuki Shimonishi, Tutomu Murase
NEC Corp.

Sep. 2, 2006 at PFLDnet2006
h-shimonishi@cd.jp.nec.com

Motivation

- Throughput degradation of TCP in high-speed and long distance environment
 - (Non-congestion) packet losses cause throughput degradation
 - Slow recovery from packet losses





New TCP variants

(1) High Speed TCP (HS-TCP) [S. Floyd, RFC 3649, 2003]

- Tune congestion window increase/decrease based on congestion window size

(2) Scalable TCP [T. Kelly, PFDnet03, 2003]

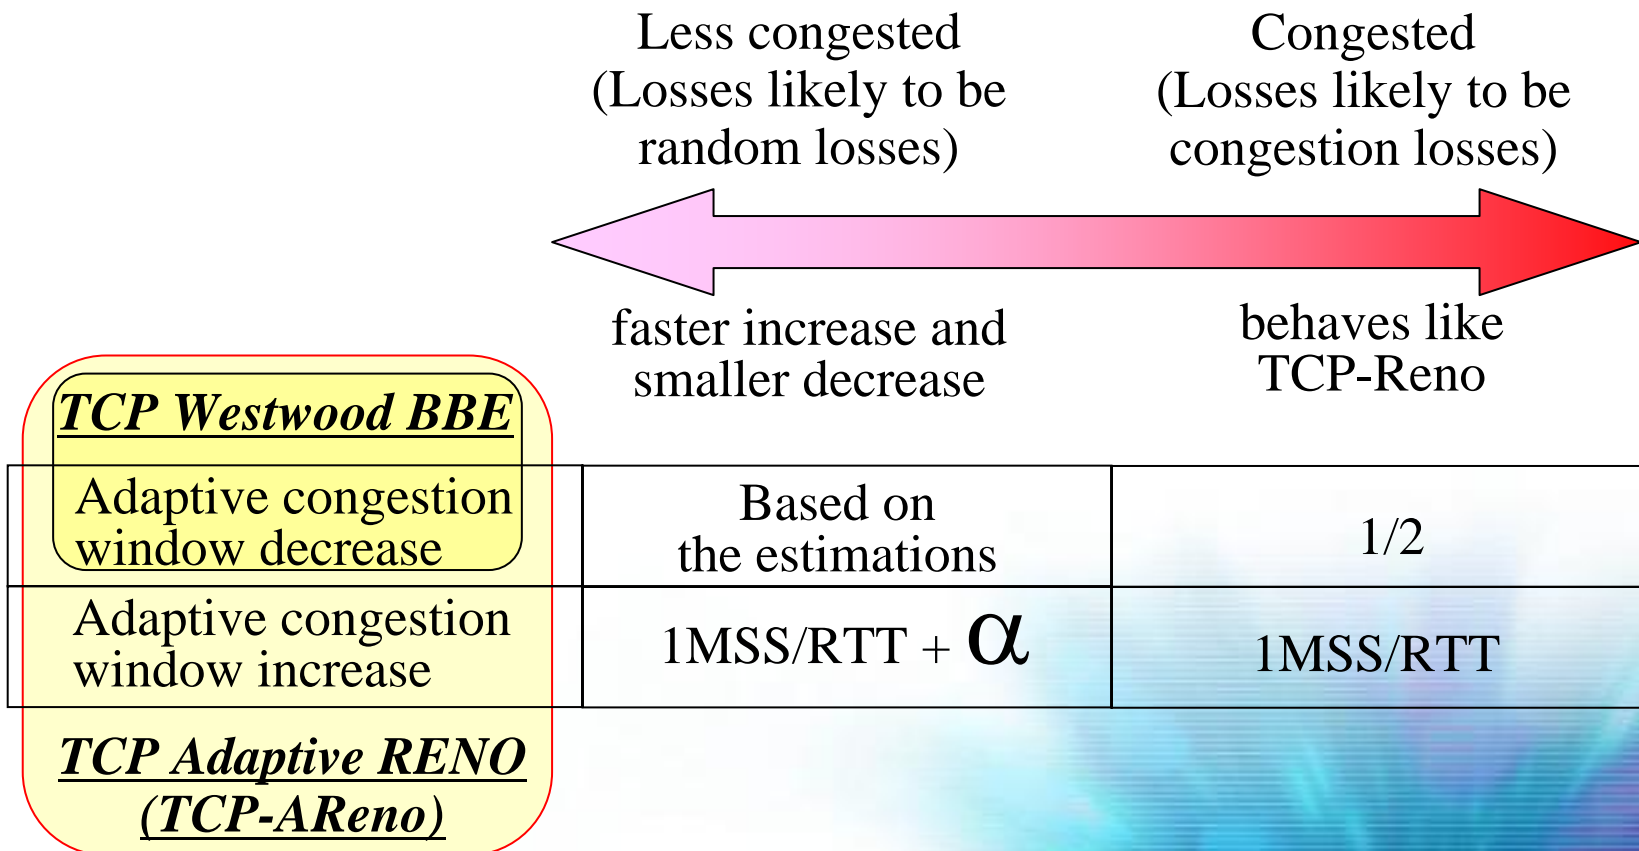
(3) FAST [S. H. Low, draft-jwl-tcp-fast-01.txt, 2003]



Challenge: friendliness to existing protocols (TCP-Reno)

Improving efficiency-friendliness tradeoff

- Tune congestion window increase/decrease based on congestion estimation

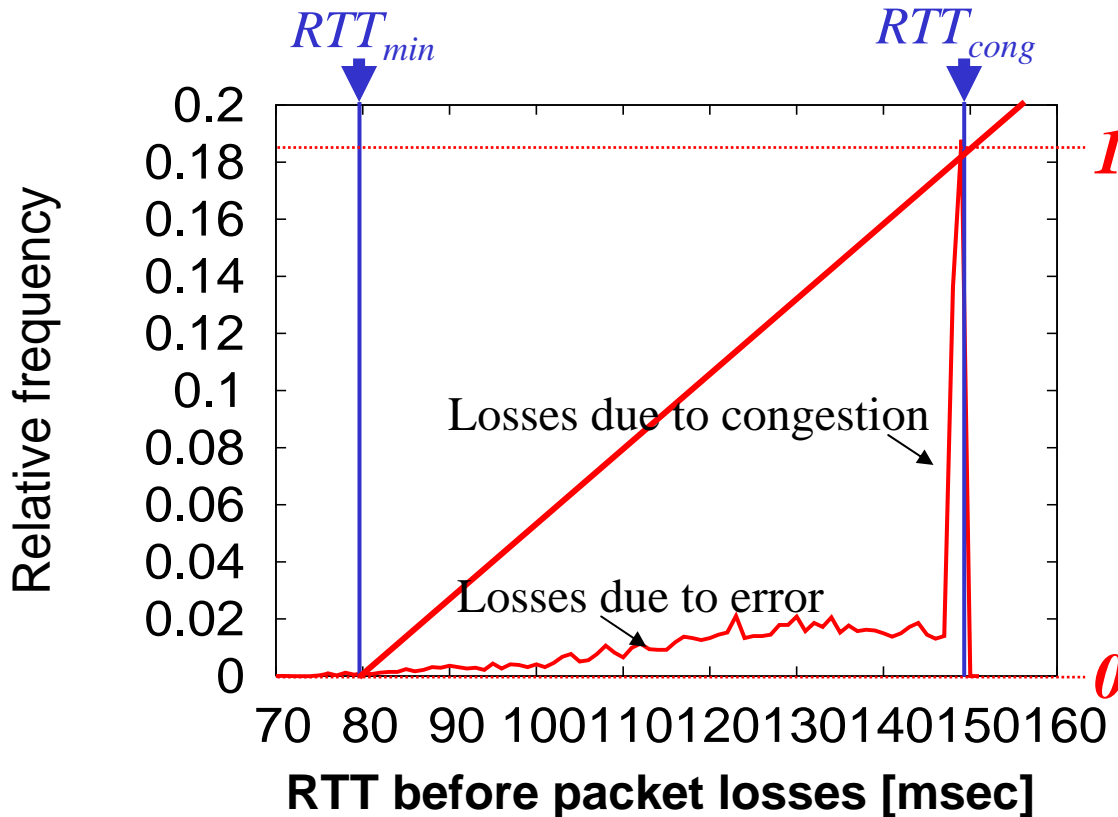


Congestion level estimation

- “Congestion” ?
 - Congestion := a moment when coexisting TCP-Reno flows likely to reduce congestion window
:= a moment when packet losses likely to happen
- Possible condition indicator
 - Network feedbacks like (multi-bit) ECN
 - One way delay / **RTT**

Congestion level estimation (cont'd)

- RTT dynamic range between RTT_{min} and RTT_{cong}
 - RTT_{min} : minimum RTT (= propagation delay)
 - RTT_{cong} : RTT value when packet losses likely to happen
- Congestion := position of current RTT in RTT dynamic range



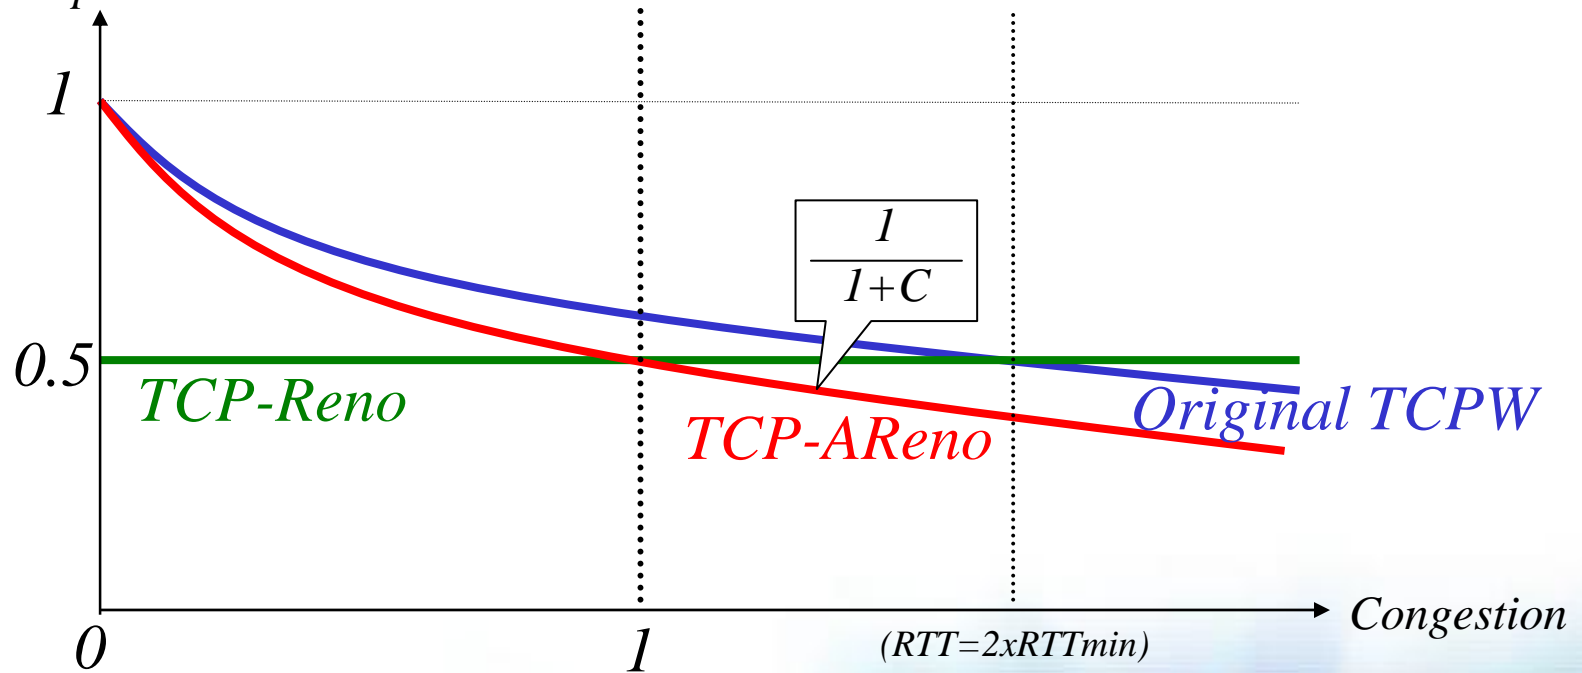
RTT_{cong} is an exponential average of RTT right before losses

Congestion level

$$c = \frac{RTT - RTT_{min}}{RTT_{cong} - RTT_{min}}$$

Congestion window reduction

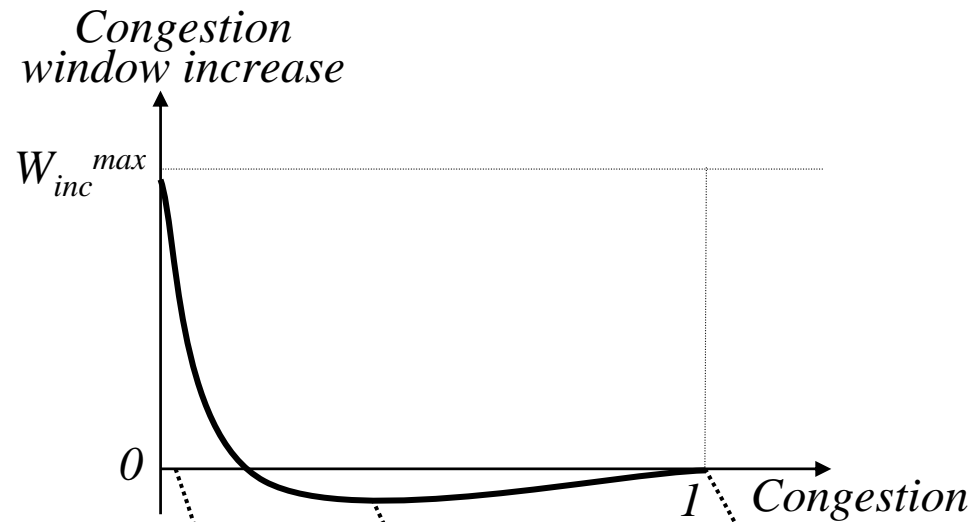

*Congestion window reduction
after a packet loss*



Should be a random loss
Smaller reduction

Should be a congestion loss
Halve the window

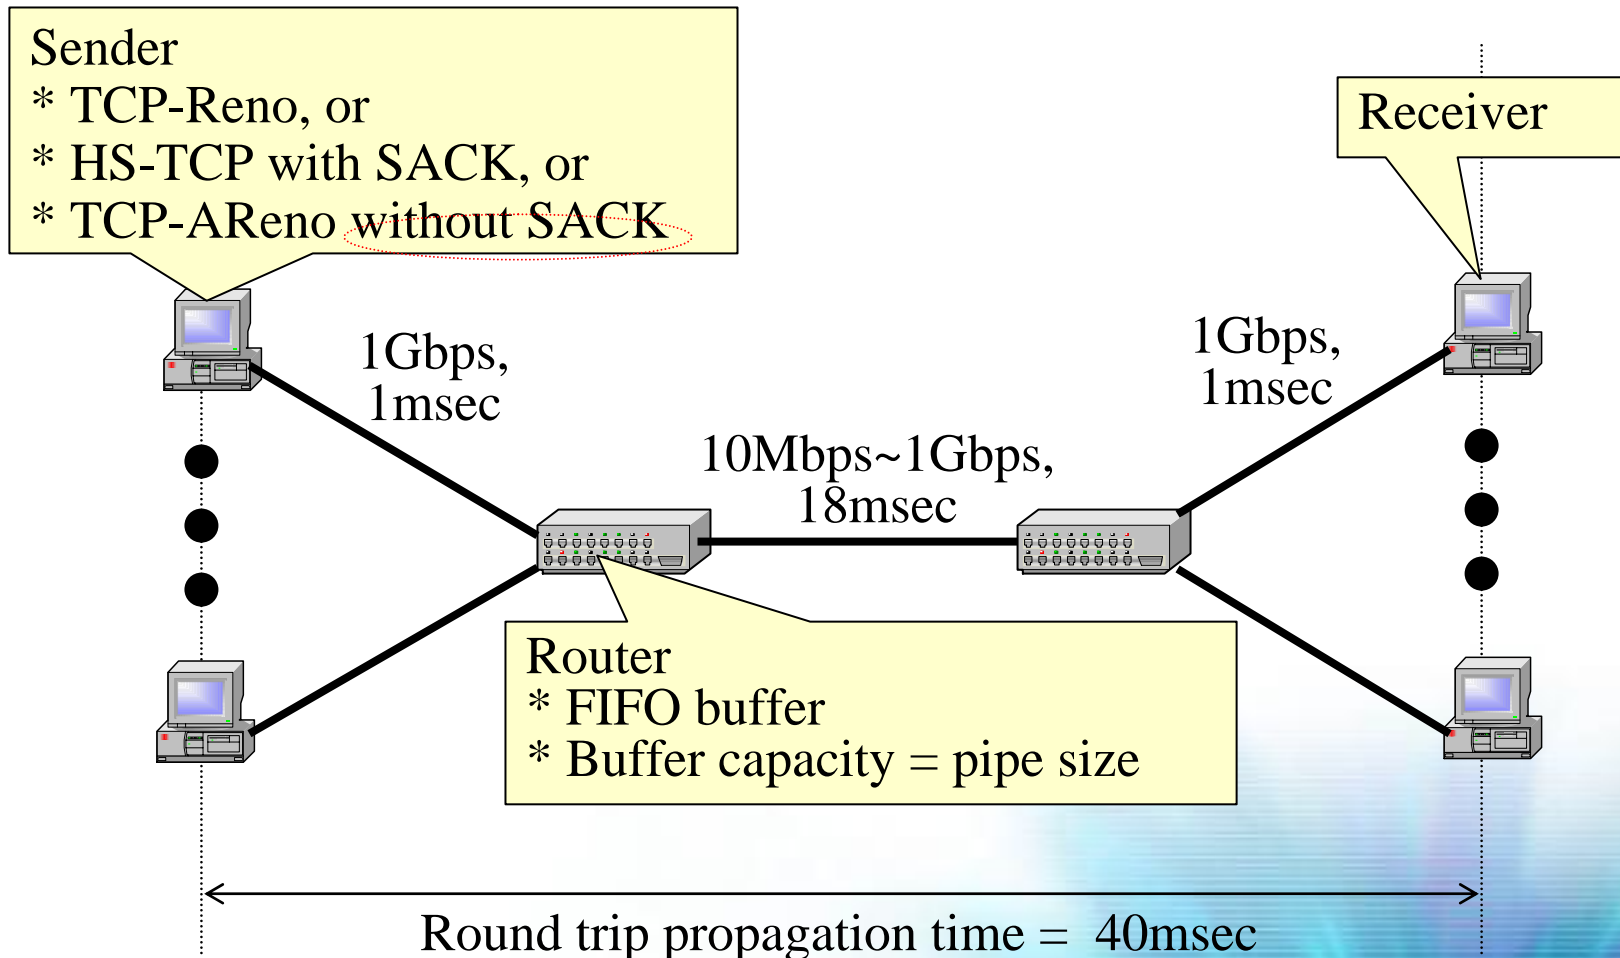
Congestion window increase



Congestion window = $W_{base} + W_{probe}$

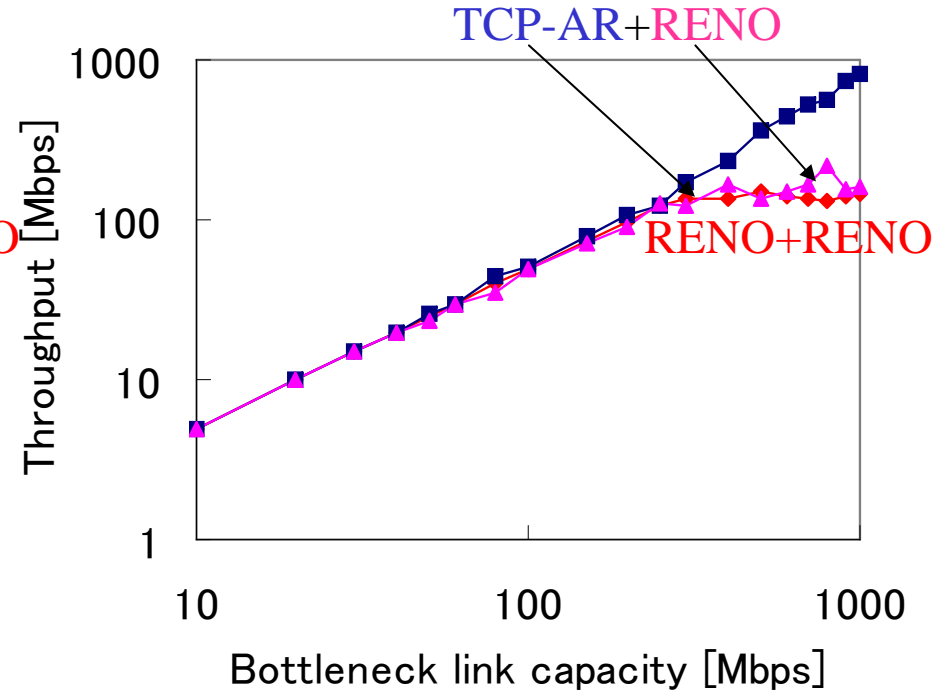

- W_{base} : increased like TCP-Reno
- W_{probe} : $c \approx 0$; increase faster – quickly fill the capacity
 $0 < c < 1$; W_{probe} converges to 0 – friendliness to TCP-Reno
 $c \approx 1$; No increase/decrease in W_{probe} part – packet loss behavior like Reno
 - avoid multiple losses
 - avoid synchronization issue

Simulation study



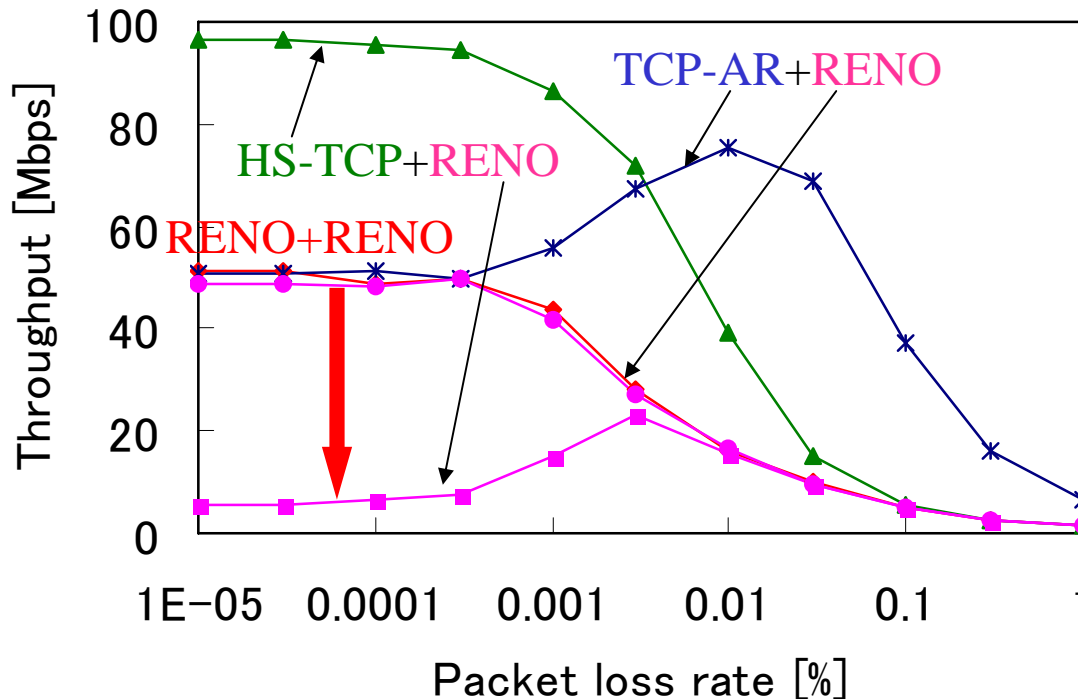
Scalability and friendliness

- Two competing flows with various link capacities
 - TCP-Reno v.s. TCP-Reno/HS-TCP/TCP-AReno
 - 10^{-6} random packet loss



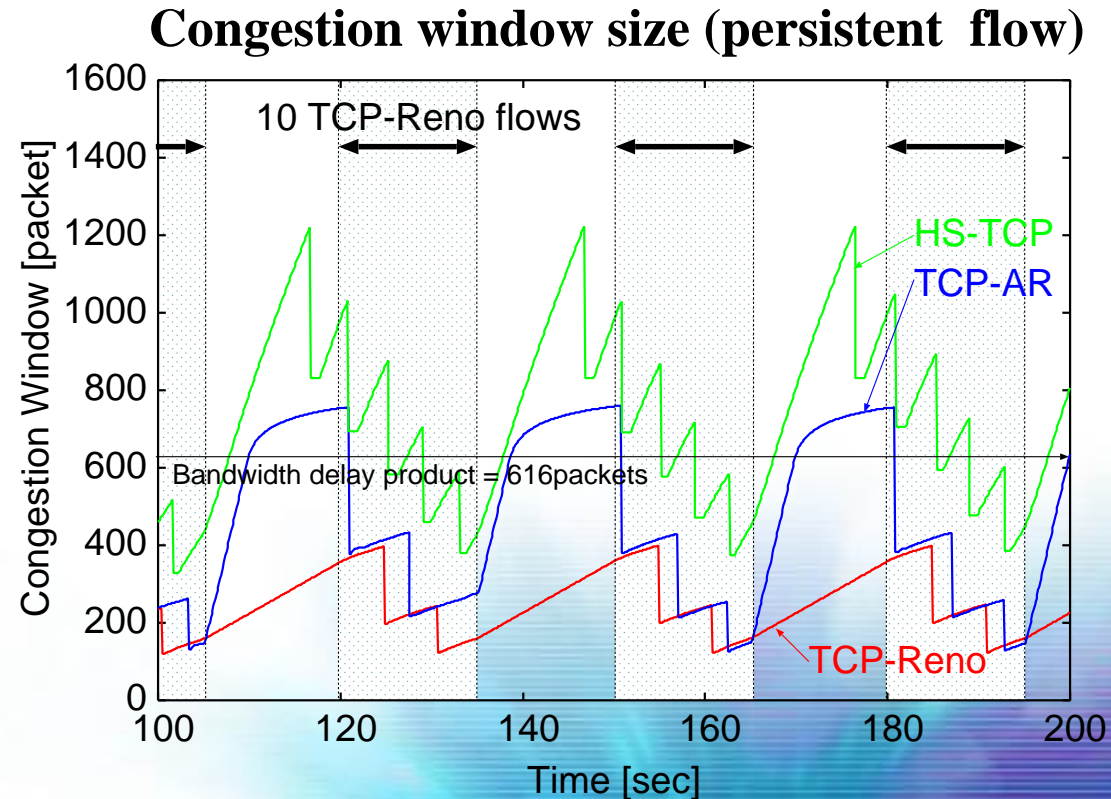
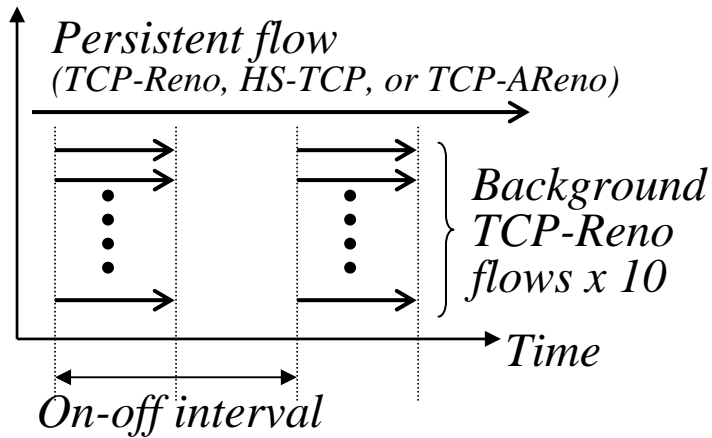
Robustness to random packet losses

- Two competing flows with various random packet loss rate
 - TCP-Reno and TCP-Reno/HS-TCP/TCP-AReno
 - 100Mbps link capacity



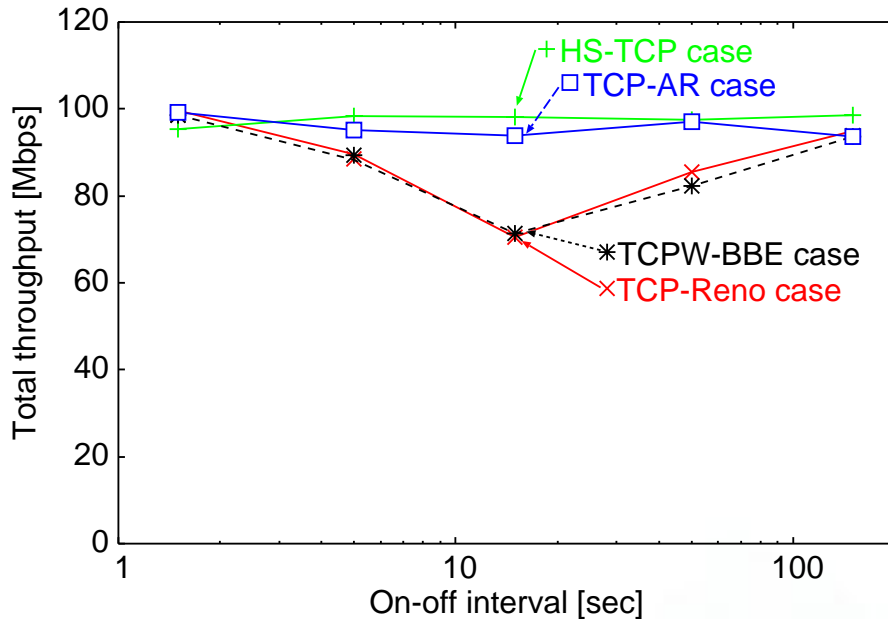
Coexisting with on-off flows (1)

- One persistent flow: TCP-Reno, or HS-TCP, or TCP-AReno
- 10 background on-off flows: TCP-Reno

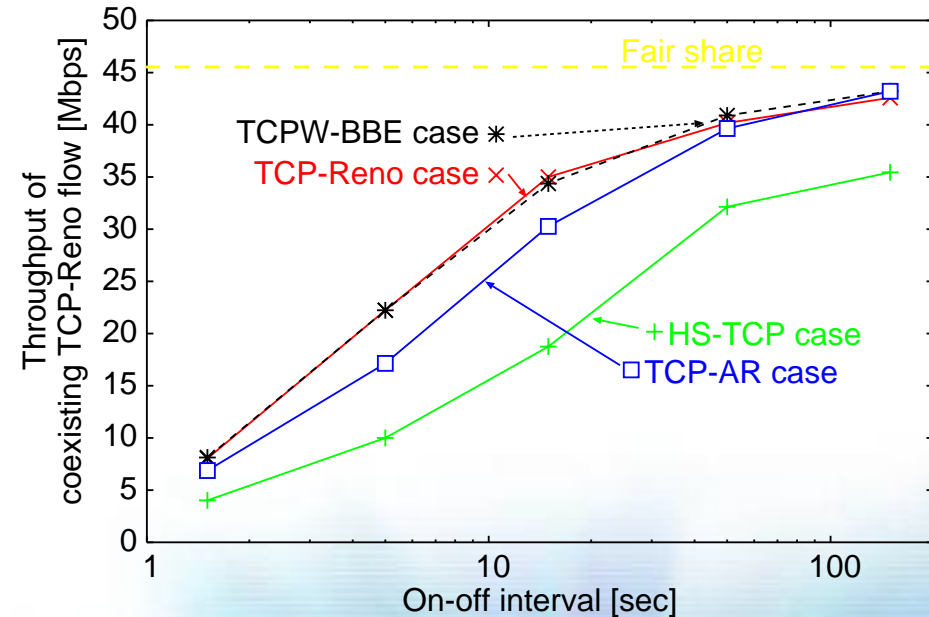


Coexisting with on-off flows (2)

Total throughput (1 persistent and 10 on/off flows)




Aggregated throughput of 10 on/off flows

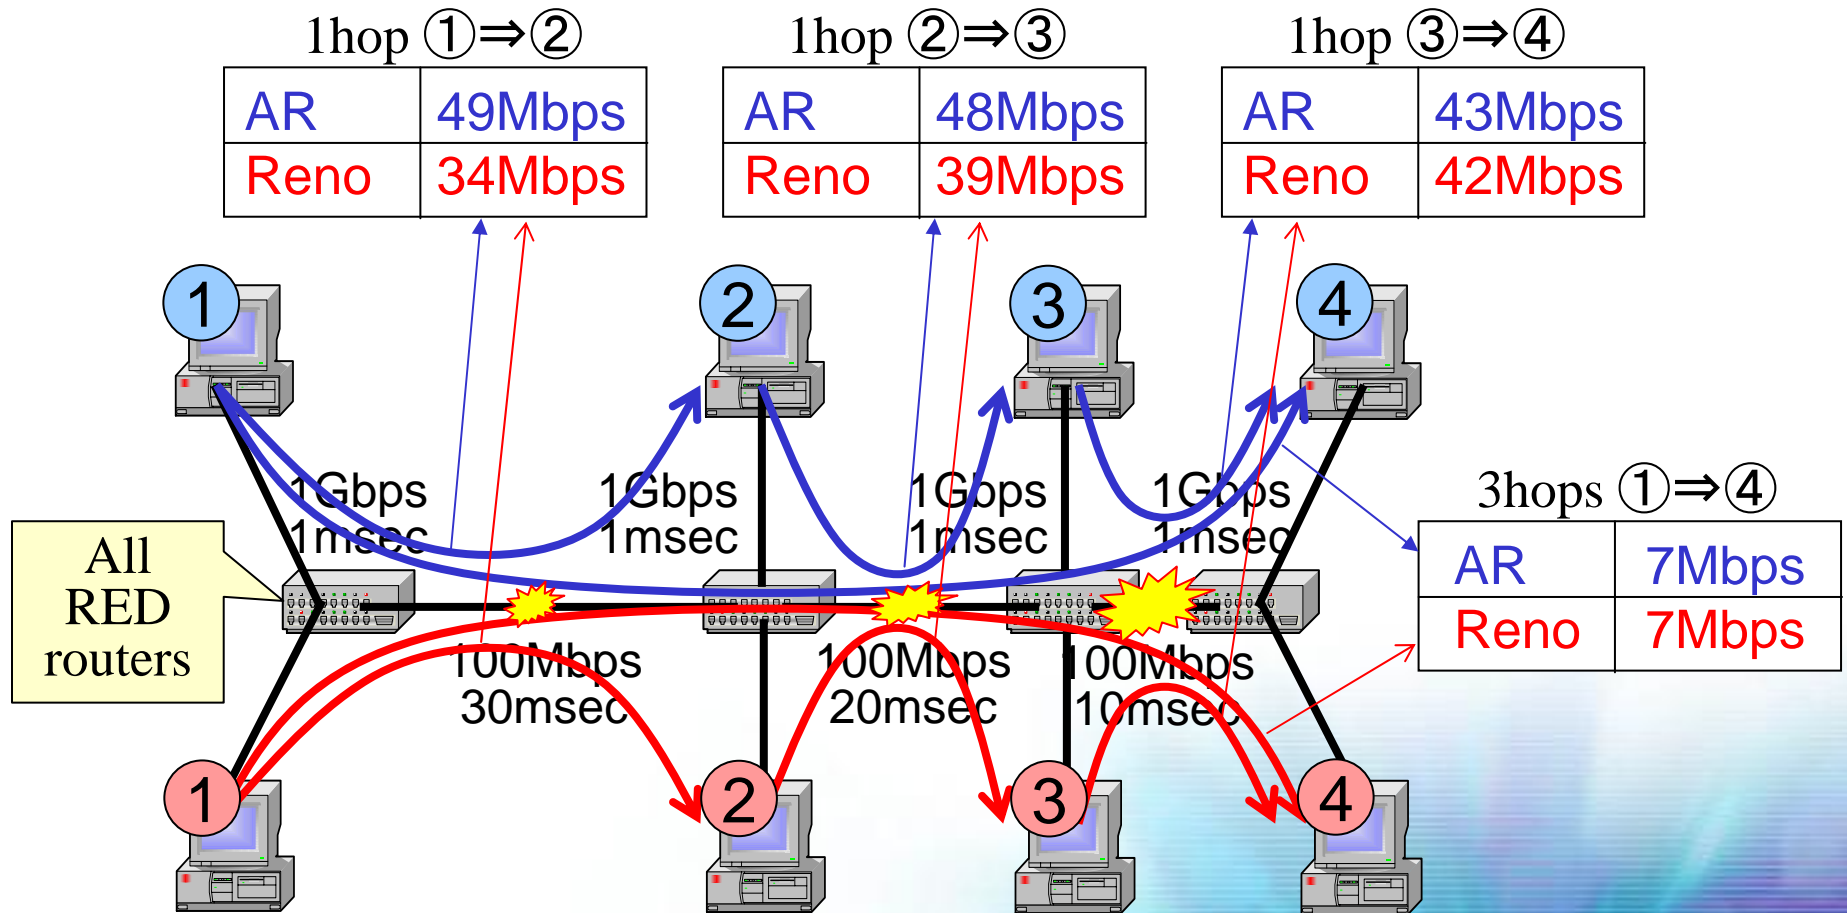


Reverse traffic case


3hops ①⇒④		2hops ②⇒④		1hop ③⇒④	
AR	20Mbps	AR	14Mbps	AR	11Mbps
Reno	24Mbps	Reno	13Mbps	Reno	12Mbps



Multiple bottleneck case



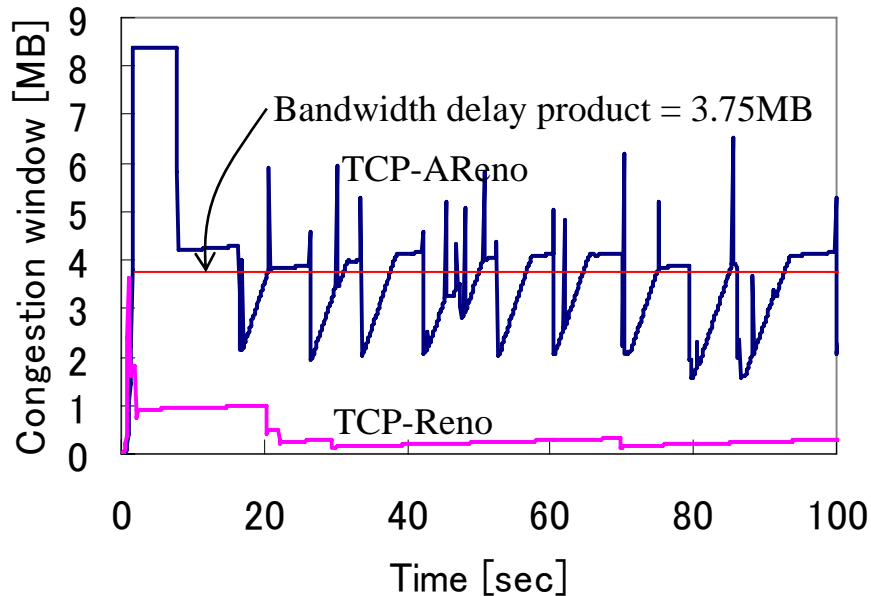
Experimental study




Laboratory measurements (1)

- 100Mbps bottleneck
- 10^{-5} random packet loss
- 100ms RTT

Separate
TCP-Reno or TCP-AReno

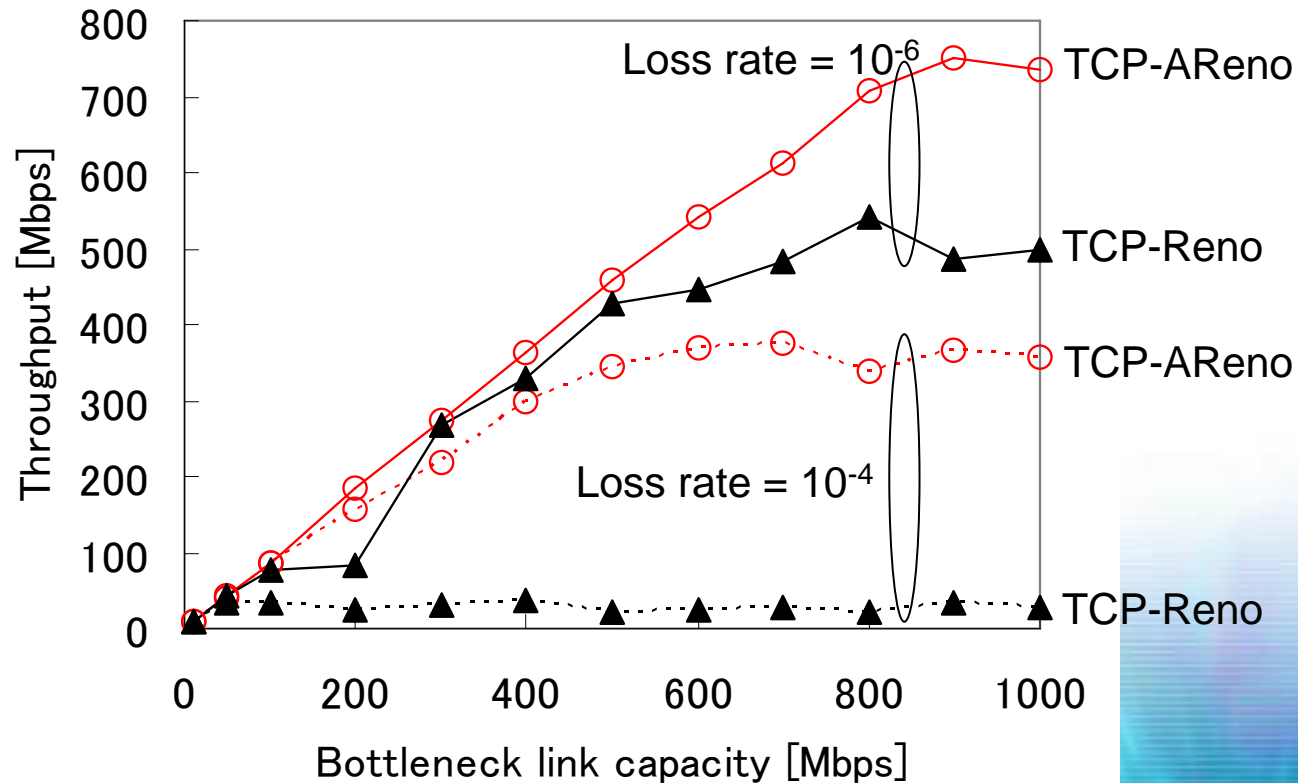


Competing
TCP-Reno v.s. TCP-AReno




Laboratory measurements (2)

- A lone TCP-Reno or TCP-AReno
- 20ms RTT



Internet measurements

- 15msec Internet line + 25msec network emulator
- No additional random packet losses at emulator



Summary

- TCP Adaptive Reno
- Rely on congestion estimation, via buffer estimation
- Whenever it finds congestion, just behaves like Reno
Otherwise, goes faster